

Database Project (Fall 2023)

HW3 (100pts, Due date: Nov 3)

Student ID: 2020315798

Student Name: Choi Jin Woo

Instruction: In this homework, we provide you with a jupyter notebook file (DBP_HW3.ipynb). You should follow the instructions in these documents carefully.

Submit two files as follows:

- DBP_HW3_StudentID.zip
- DBP_HW3_StudentID.ipynb
- DBP_HW3_StudentID.pdf

[Spark SQL]

1. [20pts] The following data is related to the manufacturing process of computer manufacturers.

Data:

- Already in the manufacturing process → used for problem (a):

```
[Row(phase='packing', model='book_pro', serial='book_pro1'),  
Row(phase='packing', model='book_pro', serial='book_pro2'),  
Row(phase='packing', model='plus', serial='plus1'),  
Row(phase='packing', model='book_pro', serial='book_pro3'),  
Row(phase='packing', model='plus', serial='plus2'),  
Row(phase='inspection', model='book_pro', serial='book_pro4'),  
Row(phase='inspection', model='plus', serial='plus3'),  
Row(phase='inspection', model='book_pro', serial='book_pro5'),  
Row(phase='inspection', model='book_pro', serial='book_pro6')]
```

- The manufacturing process to add → used for problem (b):

```
[Row(phase='assembly', model='book_pro', serial='book_pro7'),  
Row(phase='assembly', model='plus', serial='plus4')]
```

(a) Create a DataFrame with the given data and display the generated DataFrame.

[Answer]

Enter your code and result here. You must show your result (captured image).

▼ (a) Create a DataFrame with the given data and display the generated DataFrame.

```

# ===== EDIT HERE =====

df = sc.parallelize([
    Row(phase='packing', model='book_pro', serial='book_pro1'),
    Row(phase='packing', model='book_pro', serial='book_pro2'),
    Row(phase='packing', model='plus', serial='plus1'),
    Row(phase='packing', model='book_pro', serial='book_pro3'),
    Row(phase='packing', model='plus', serial='plus2'),
    Row(phase='inspection', model='book_pro', serial='book_pro4'),
    Row(phase='inspection', model='plus', serial='plus3'),
    Row(phase='inspection', model='book_pro', serial='book_pro5'),
    Row(phase='inspection', model='book_pro', serial='book_pro6')
]).toDF()

# =====

df.show(truncate=False)

```

phase	model	serial
packing	book_pro	book_pro1
packing	book_pro	book_pro2
packing	plus	plus1
packing	book_pro	book_pro3
packing	plus	plus2
inspection	book_pro	book_pro4
inspection	plus	plus3
inspection	book_pro	book_pro5
inspection	book_pro	book_pro6

(b) After adding *two laptops to the manufacturing process*, find the number of products for each model.

[Answer]

Enter your code and result here. You must show your result (captured image).

▼ (b) After adding two laptops to the manufacturing process, find the number of products for each model.

```

# ===== EDIT HERE =====

df2 = sc.parallelize([
    Row(phase='assembly', model='book_pro', serial='book_pro7'),
    Row(phase='assembly', model='plus', serial='plus4')
]).toDF()

new_df = df.unionByName(df2).groupBy("model").count()

# =====

new_df.show(truncate=False)

```

model	count
book_pro7	1
plus	4

(c) Group the data in the joined DataFrame by 'phases' and count the number of data for each phase.

[Answer] Enter your code and result here. You have to show your snapshot result.

▼ (c) Group the data in the joined DataFrame by 'phases' and count the number of data for each phase.

```

# ===== EDIT HERE =====

group_df = df.unionByName(df2).groupBy("phase").count()

# =====

group_df.show(truncate=False)

```

phase	count
packing	5
inspection	4
assembly	2

2. [20pts] The following data are *manufacturing process* and *customer information* for computer manufacturers.

- Data 1(manufacturing process):

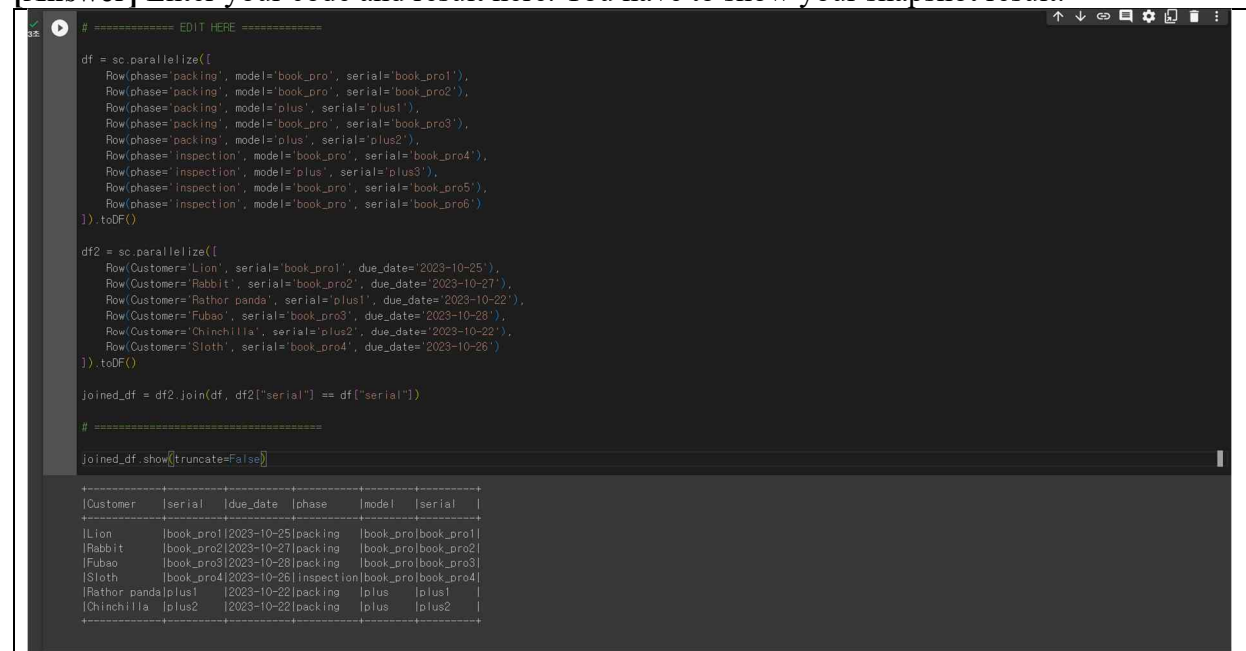
```
[Row(phase='packing', model='book_pro', serial='book_pro1'),
Row(phase='packing', model='book_pro', serial='book_pro2'),
Row(phase='packing', model='plus', serial='plus1'),
Row(phase='packing', model='book_pro', serial='book_pro3'),
Row(phase='packing', model='plus', serial='plus2'),
Row(phase='inspection', model='book_pro', serial='book_pro4'),
Row(phase='inspection', model='plus', serial='plus3'),
Row(phase='inspection', model='book_pro', serial='book_pro5'),
Row(phase='inspection', model='book_pro', serial='book_pro6')]
```

- Data 2(customer information):

```
[Row(Customer='Lion', serial='book_pro1', due_date='2023-10-25'),
Row(Customer='Rabbit', serial='book_pro2', due_date='2023-10-27'),
Row(Customer='Rathor panda', serial='plus1', due_date='2023-10-22'),
Row(Customer='Fubao', serial='book_pro3', due_date='2023-10-28'),
Row(Customer='Chinchilla', serial='plus2', due_date='2023-10-22'),
Row(Customer='Sloth', serial='book_pro4', due_date='2023-10-26')]
```

- (a) Create a DataFrame for the two given data and join Data 1 with Data 2 using an inner join based on the 'serial' column. (left side: Data 2, right side: Data 1)

[Answer] Enter your code and result here. You have to show your snapshot result.



```
# ===== EDIT HERE =====

df = sc.parallelize([
    Row(phase='packing', model='book_pro', serial='book_pro1'),
    Row(phase='packing', model='book_pro', serial='book_pro2'),
    Row(phase='packing', model='plus', serial='plus1'),
    Row(phase='packing', model='book_pro', serial='book_pro3'),
    Row(phase='packing', model='plus', serial='plus2'),
    Row(phase='inspection', model='book_pro', serial='book_pro4'),
    Row(phase='inspection', model='plus', serial='plus3'),
    Row(phase='inspection', model='book_pro', serial='book_pro5'),
    Row(phase='inspection', model='book_pro', serial='book_pro6')
]).toDF()

df2 = sc.parallelize([
    Row(Customer='Lion', serial='book_pro1', due_date='2023-10-25'),
    Row(Customer='Rabbit', serial='book_pro2', due_date='2023-10-27'),
    Row(Customer='Rathor panda', serial='plus1', due_date='2023-10-22'),
    Row(Customer='Fubao', serial='book_pro3', due_date='2023-10-28'),
    Row(Customer='Chinchilla', serial='plus2', due_date='2023-10-22'),
    Row(Customer='Sloth', serial='book_pro4', due_date='2023-10-26')
]).toDF()

joined_df = df2.join(df, df2["serial"] == df["serial"])

# =====

joined_df.show(truncate=False)
```

Customer	serial	due_date	phase	model	serial
Lion	book_pro1	2023-10-25	packing	book_pro	book_pro1
Rabbit	book_pro2	2023-10-27	packing	book_pro	book_pro2
Fubao	book_pro3	2023-10-28	packing	book_pro	book_pro3
Sloth	book_pro4	2023-10-26	inspection	book_pro	book_pro4
Rathor panda	plus1	2023-10-22	packing	plus	plus1
Chinchilla	plus2	2023-10-22	packing	plus	plus2

- (b) Use an SQL query to select the data from the joined DataFrame where the 'due_date' is on or after '2023-10-25'. And briefly explain the method you used.

[Answer] Enter your code and result here. You have to show your snapshot result.

(b) Use an SQL query to select the data from the joined DataFrame where the 'due_date' is on or after '2023-10-25'. And briefly explain the method you used.

```

===== EDIT HERE =====
joined_df.createOrReplaceTempView("joined_df")

sql_query = """
SELECT *
FROM joined_df
WHERE due_date >= '2023-10-25'
"""

sql_df = spark.sql(sql_query)

# =====
sql_df.show(truncate=False)

```

Customer	serial	due_date	phase	model	serial
Llion	book_pro1	2023-10-25	packing	book_pro1	book_pro1
Rabbit	book_pro2	2023-10-27	packing	book_pro1	book_pro2
Fubao	book_pro3	2023-10-28	packing	book_pro1	book_pro3
Sloth	book_pro4	2023-10-26	inspection	book_pro1	book_pro4

[createOrReplaceTempView]
To save DataFrame into temporary SQL table named “joined_df”

[sql_query]
SQL query that filters where “due_date” is on or after “2023-10-25”

[spark.sql()]
Returns new DataFrame with query result

[Spark ML]

3. [60pts] We provide you with a *Fashion-MNIST* dataset.

Dataset Description:

Training set: 60,000 examples

Test set: 10,000 examples

Each example is a 28x28 grayscale image associated with a label from 10 classes.

Label	Description
0	T-shirt/top
1	Trouser
2	Pullover
3	Dress
4	Coat
5	Sandal
6	Shirt
7	Sneaker
8	Bag
9	Ankle boot

For more information, visit this website: <https://github.com/zalandoresearch/fashion-mnist>

(a) Load the provided dataset, convert it into a DataFrame, and show it. You should follow the following instructions.

-Instructions 1: Assemble the features into a vector column and name the column “features.”

-Instructions 2: Rename the target column to “label.”

[Answer] Enter your code and result here. You have to show your snapshot result.

```
(a) Load the provided dataset, convert it into a DataFrame, and show it.

[30] def create_dataframe(data_path):

    fashion_mnist_data = spark.read.option('header', 'true').option('inferSchema', 'true').csv(data_path)

    col_list = ['pixel'+str(i+1) for i in range(784)]

    # ===== EDIT HERE =====

    # Instruction 1: Assemble the features into a vector column and name the column "features"
    assembler = VectorAssembler(inputCols=col_list, outputCol="features")
    fashion_mnist_data = assembler.transform(fashion_mnist_data)

    # Instruction 2: Rename the target column to "label"
    fashion_mnist_data = fashion_mnist_data.withColumnRenamed("digit", "label")

    # ===== EDIT HERE =====

    return fashion_mnist_data

# ===== EDIT HERE =====
# OPT/NOTE: You can also change the data path.
train_data_path = "fashion-mnist-train.csv"
test_data_path = "fashion-mnist-test.csv"
# =====
train_data = create_dataframe(train_data_path)
test_data = create_dataframe(test_data_path)

train_data.show()
test_data.show()
```

only showing top 20 rows

(b) Train models to classify the classes of the Fashion MNIST dataset and report the results for test data. The models used are **Logistic Regression, Decision Tree, and Random Forest**.

For detailed explanations of the usage of each model, please refer to the official documentation below.

- Logistic Regression: [\[Link\]](#)
- Decision Tree: [\[Link\]](#)
- Random Forest: [\[Link\]](#)

[Answer] Fill in the table below.

	Logistic Regression	Decision Tree	Random Forest
Test accuracy	0.7561	0.7441	0.7724

[Answer] Enter your code and result here. You have to show your snapshot result.

Logistic Regression

Reference: <https://spark.apache.org/docs/latest/api/python/reference/api/pyspark.ml.classification.LogisticRegression.html>

```
from pyspark.ml.classification import LogisticRegression

# Training and Test
# ===== EDIT HERE =====
lr = LogisticRegression(maxIter=100, regParam=0.05, elasticNetParam=0.5)
lr_model = lr.fit(train_data)
lr_preds = lr_model.transform(test_data)
lr_accuracy = evaluator.evaluate(lr_preds)

# =====

print(f"Accuracy: {lr_accuracy}")

Accuracy: 0.7560803665844201
```

Decision Tree

Reference: <https://spark.apache.org/docs/latest/api/python/reference/api/pyspark.ml.classification.DecisionTreeClassifier.html>

```
from pyspark.ml.classification import DecisionTreeClassifier

# On model declaration, fix seed, maxDepth, to the following values
seed = 2023

# ===== EDIT HERE =====
dt = DecisionTreeClassifier(seed=seed, maxDepth=15)
dt_model = dt.fit(train_data)
dt_preds = dt_model.transform(test_data)
dt_accuracy = evaluator.evaluate(dt_preds)

# =====

print(f"Accuracy: {dt_accuracy}")

Accuracy: 0.7440958759252732
```

Random Forest

Reference: <https://spark.apache.org/docs/latest/api/python/reference/api/pyspark.ml.classification.RandomForestClassifier.html>

```
[12] from pyspark.ml.classification import RandomForestClassifier

# On model declaration, fix seed to the following values
seed = 2023

# ===== EDIT HERE =====
rfc = RandomForestClassifier(seed=seed)
rfc_model = rfc.fit(train_data)
rfc_preds = rfc_model.transform(test_data)
rfc_accuracy = evaluator.evaluate(rfc_preds)

# =====

print(f"Accuracy: {rfc_accuracy}")

Accuracy: 0.7724121724826695
```